

Assignment 4 Project Report

Li Jiaqi 120090727

Overview

In this assignment, I have completed both the basic single-directory File System and the bonus task of tree-structured directory File System. The following content will describe the relevant information of these two tasks I completed.

Environment

OS Version

I use the university's High Performance Cluster (HPC) for testing and running the CUDA program. The nodes run on a CentOS version 7.5.1804.

```
[120090727@node21 ~]$ cat /etc/redhat-release  
CentOS Linux release 7.5.1804 (Core)
```

Kernel version

This is the kernel version of the HPC. Other versions should also be OK.

```
[120090727@node21 ~]$ uname -r  
3.10.0-862.el7.x86_64
```

CUDA Version

I use the CUDA compiler version 11.7 for compiling the CUDA program.

```
[120090727@node21 ~]$ nvcc --version
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2022 NVIDIA Corporation
Built on Wed_Jun__8_16:49:14_PDT_2022
Cuda compilation tools, release 11.7, V11.7.99
Build cuda_11.7.r11.7/compiler.31442593_0
```

GPU Info

For each node in the HPC, it is equipped with a Quadro RTX 4000 GPU. Each time the program only runs on one allocated node. I have also tested my programs on a RTX3090 GPU.

NVIDIA-SMI 515.65.01				Driver Version: 515.65.01		CUDA Version: 11.7	
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr. ECC	
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute M.	MIG M.
=====							
0	Quadro RTX 4000	Off	00000000:AF:00.0	Off			N/A
36%	63C	P0	59W / 125W	250MiB / 8192MiB	100%	Default	N/A
=====							

Running the program

Basic task compilation and running

To compile: inside the `source/` folder, there is a file named `slurm.sh`. On the HPC with slurm installed, we can directly use the shell script to compile and run the executable:

```
sbatch slurm.sh
```

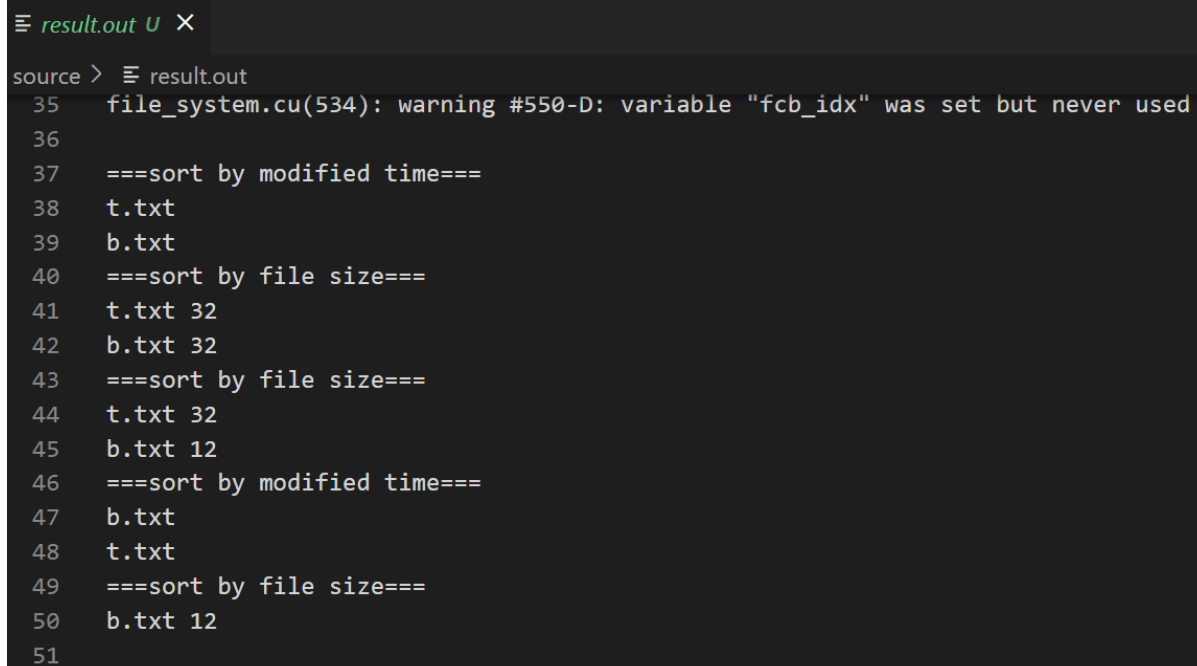
On a device without slurm, one can first compile using

```
nvcc --relocatable-device-code=true main.cu user_program.cu file_system.cu -o test
```

and run `./test` to run the program (might need `srun` in the cluster).

Basic task sample outputs

On the first test program: (The first 36 lines are compiler warnings)



```
source > result.out
35 file_system.cu(534): warning #550-D: variable "fcb_idx" was set but never used
36
37 ===sort by modified time===
38 t.txt
39 b.txt
40 ===sort by file size===
41 t.txt 32
42 b.txt 32
43 ===sort by file size===
44 t.txt 32
45 b.txt 12
46 ===sort by modified time===
47 b.txt
48 t.txt
49 ===sort by file size===
50 b.txt 12
51
```

On the second test program: (The first 36 lines are compiler warnings)

```

source > ≡ result.out

37  ===sort by modified time===
38  t.txt
39  b.txt
40  ===sort by file size===
41  t.txt 32
42  b.txt 32
43  ===sort by file size===
44  t.txt 32
45  b.txt 12
46  ===sort by modified time===
47  b.txt
48  t.txt
49  ===sort by file size===
50  b.txt 12
51  ===sort by file size===
52  *ABCDEFGHIJKLMNOPQR 33
53  )ABCDEFGHIJKLMNOPQR 32
54  (ABCDEFGHIJKLMNOPQR 31
55  'ABCDEFGHIJKLMNOPQR 30
56  &ABCDEFGHIJKLMNOPQR 29
57  %ABCDEFGHIJKLMNOPQR 28
58  $ABCDEFGHIJKLMNOPQR 27
59  #ABCDEFGHIJKLMNOPQR 26
60  "ABCDEFGHIJKLMNOPQR 25
61  !ABCDEFGHIJKLMNOPQR 24
62  b.txt 12
63  ===sort by modified time===
64  *ABCDEFGHIJKLMNOPQR
65  )ABCDEFGHIJKLMNOPQR
66  (ABCDEFGHIJKLMNOPQR
67  'ABCDEFGHIJKLMNOPQR
68  &ABCDEFGHIJKLMNOPQR
69  b.txt
70

```

On the third test program: (there are too many lines and we only display the begin, middle and end)

```

41  ===sort by modified time===
42  t.txt
43  b.txt
44  ===sort by file size===
45  t.txt 32
46  b.txt 32
47  ===sort by file size===
48  t.txt 32
49  b.txt 12
50  ===sort by modified time===
51  b.txt
52  t.txt
53  ===sort by file size===
54  b.txt 12
55  ===sort by file size===
56  *ABCDEFGHJKLMNOPQR 33
57  )ABCDEFGHJKLMNOPQR 32
58  (ABCDEFGHJKLMNOPQR 31
59  'ABCDEFGHJKLMNOPQR 30
60  &ABCDEFGHJKLMNOPQR 29
61  %ABCDEFGHJKLMNOPQR 28
62  $ABCDEFGHJKLMNOPQR 27
63  #ABCDEFGHJKLMNOPQR 26
64  "ABCDEFGHJKLMNOPQR 25
65  !ABCDEFGHJKLMNOPQR 24
66  b.txt 12
67  ===sort by modified time===
68  *ABCDEFGHJKLMNOPQR
69  )ABCDEFGHJKLMNOPQR
70  (ABCDEFGHJKLMNOPQR
71  'ABCDEFGHJKLMNOPQR
72  &ABCDEFGHJKLMNOPQR
73  b.txt
74  ===sort by file size===
75  ~ABCDEFGHJKLM 1024
76  }ABCDEFGHJKLM 1023
77  |ABCDEFGHJKLM 1022
78  {ABCDEFGHJKLM 1021
79  zABCDEFGHJKLM 1020
80  yABCDEFGHJKLM 1019
81  xABCDEFGHJKLM 1018
82  wABCDEFGHJKLM 1017
83  vABCDEFGHJKLM 1016
84  uABCDEFGHJKLM 1015
85  tABCDEFGHJKLM 1014

```

```
1057    DA 42
1058    CA 41
1059    BA 40
1060    AA 39
1061    @A 38
1062    ?A 37
1063    >A 36
1064    =A 35
1065    <A 34
1066    *ABCDEFGHJKLMNOPQR 33
1067    ;A 33
1068    )ABCDEFGHJKLMNOPQR 32
1069    :A 32
1070    (ABCDEFGHJKLMNOPQR 31
1071    9A 31
1072    'ABCDEFGHJKLMNOPQR 30
1073    8A 30
1074    &ABCDEFGHJKLMNOPQR 29
1075    7A 29
1076    6A 28
1077    5A 27
1078    4A 26
1079    3A 25
1080    2A 24
1081    b.txt 12
1082    ===sort by file size===
1083    EA 1024
1084    ~ABCDEFGHIJKLM 1024
1085    aa 1024
1086    bb 1024
1087    cc 1024
1088    dd 1024
1089    ee 1024
1090    ff 1024
1091    gg 1024
1092    hh 1024
1093    ii 1024
1094    jj 1024
```

```
2063 XA 62
2064 WA 61
2065 VA 60
2066 UA 59
2067 TA 58
2068 SA 57
2069 RA 56
2070 QA 55
2071 PA 54
2072 OA 53
2073 NA 52
2074 MA 51
2075 LA 50
2076 KA 49
2077 JA 48
2078 IA 47
2079 HA 46
2080 GA 45
2081 FA 44
2082 DA 42
2083 CA 41
2084 BA 40
2085 AA 39
2086 @A 38
2087 ?A 37
2088 >A 36
2089 =A 35
2090 <A 34
2091 *ABCDEFGHIJKLMNOPQR 33
2092 ;A 33
2093 )ABCDEFGHIJKLMNOPQR 32
2094 :A 32
2095 (ABCDEFGHIJKLMNOPQR 31
2096 9A 31
2097 'ABCDEFGHIJKLMNOPQR 30
2098 8A 30
2099 &ABCDEFGHIJKLMNOPQR 29
2100 7A 29
2101 6A 28
2102 5A 27
2103 4A 26
2104 3A 25
2105 2A 24
2106 b.txt 12
2107
```

One the fourth test case: (there are too many lines and we only screenshot the beginning and end)


```
source > ≡ result.out
25   triggering gc
26   ===sort by modified time===
27   1024-block-1023
28   1024-block-1022
29   1024-block-1021
30   1024-block-1020
31   1024-block-1019
32   1024-block-1018
33   1024-block-1017
34   1024-block-1016
35   1024-block-1015
36   1024-block-1014
37   1024-block-1013
38   1024-block-1012
39   1024-block-1011
40   1024-block-1010
41   1024-block-1009
42   1024-block-1008
43   1024-block-1007
44   1024-block-1006
45   1024-block-1005
46   1024-block-1004
47   1024-block-1003
48   1024-block-1002
49   1024-block-1001
50   1024-block-1000
51   1024-block-0999
52   1024-block-0998
53   1024-block-0997
54   1024-block-0996
55   1024-block-0995
56   1024-block-0994
57   1024-block-0993
58   1024-block-0992
59   1024-block-0991
60   1024-block-0990
61   1024-block-0989
62   1024-block-0988
63   1024-block-0987
64   1024-block-0986
65   1024-block-0985
66   1024-block-0984
67   1024-block-0983
68   1024-block-0982
69   1024-block-0981
```

source >	≡ result.out
1009	1024-block-0041
1010	1024-block-0040
1011	1024-block-0039
1012	1024-block-0038
1013	1024-block-0037
1014	1024-block-0036
1015	1024-block-0035
1016	1024-block-0034
1017	1024-block-0033
1018	1024-block-0032
1019	1024-block-0031
1020	1024-block-0030
1021	1024-block-0029
1022	1024-block-0028
1023	1024-block-0027
1024	1024-block-0026
1025	1024-block-0025
1026	1024-block-0024
1027	1024-block-0023
1028	1024-block-0022
1029	1024-block-0021
1030	1024-block-0020
1031	1024-block-0019
1032	1024-block-0018
1033	1024-block-0017
1034	1024-block-0016
1035	1024-block-0015
1036	1024-block-0014
1037	1024-block-0013
1038	1024-block-0012
1039	1024-block-0011
1040	1024-block-0010
1041	1024-block-0009
1042	1024-block-0008
1043	1024-block-0007
1044	1024-block-0006
1045	1024-block-0005
1046	1024-block-0004
1047	1024-block-0003
1048	1024-block-0002
1049	1024-block-0001
1050	1024-block-0000
1051	

Bonus Task Compilation and running

Because the bonus task shares the same template structure with the basic task, the compilation and running steps are exactly the same as above, which means we could use:

```
sbatch slurm.sh
```

Bonus task sample output

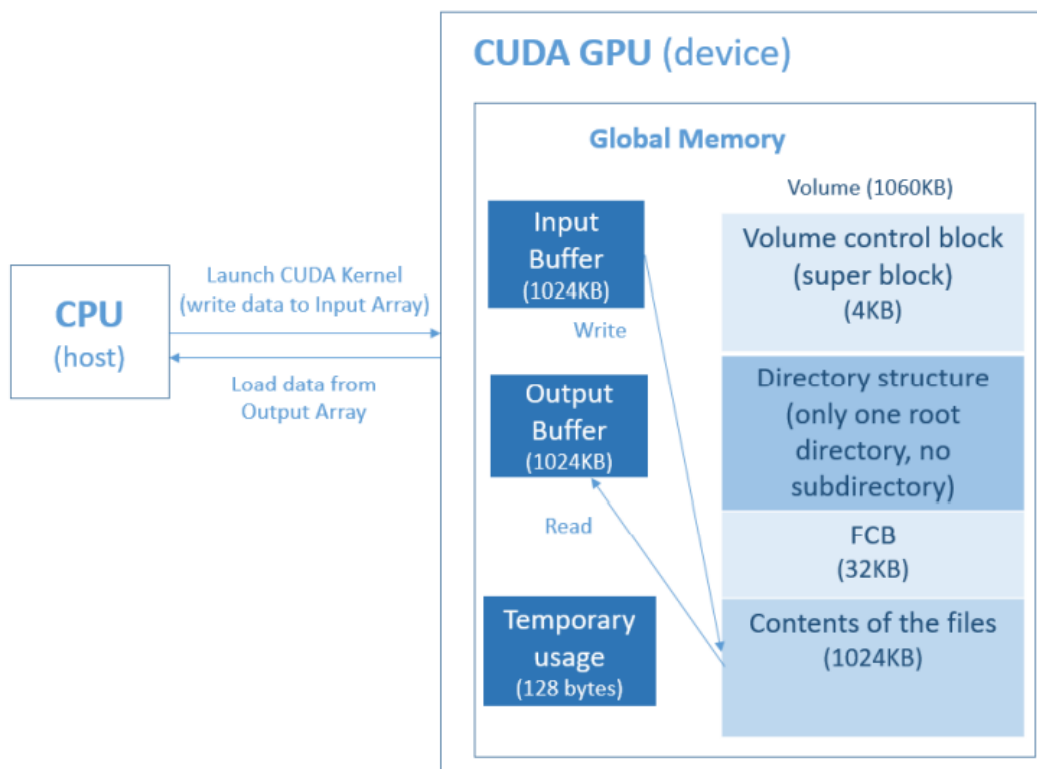
The first 91 lines are compiler warnings.

```
bonus > ≡ result.out
92  ===sort by modified time===
93  t.txt
94  b.txt
95  ===sort by file size===
96  t.txt 32
97  b.txt 32
98  ===sort by modified time===
99  app d
100 t.txt
101 b.txt
102 ===sort by file size===
103 t.txt 32
104 b.txt 32
105 app 0 d
106 ===sort by file size===
107 ===sort by file size===
108 a.txt 64
109 b.txt 32
110 soft 0 d
111 ===sort by modified time===
112 soft d
113 b.txt
114 a.txt
115 /app/soft
116 ===sort by file size===
117 B.txt 1024
118 C.txt 1024
119 D.txt 1024
120 A.txt 64
121 ===sort by file size===
122 a.txt 64
123 b.txt 32
124 soft 24 d
125 /app
126 ===sort by file size===
127 t.txt 32
128 b.txt 32
129 app 17 d
130 ===sort by file size===
131 a.txt 64
132 b.txt 32
133 ===sort by file size===
134 t.txt 32
135 b.txt 32
136 app 12 d
```

Program Design

Basic task design

In this CUDA program, we implement a single-directory file system using a limited GPU memory pool. The memory usage strictly obeys the one in the instruction, that no extra global memory is maintained or used. Temporary usage for function stack is limited.



For the task, we allocate a volume of 1060kb with the 4kb volume control block using bitmap, 32kb for 1024 FCBs, each FCB is 32 bytes. The content of files use 1024kb, divided into storage blocks each 32 bytes.

FCB Structure

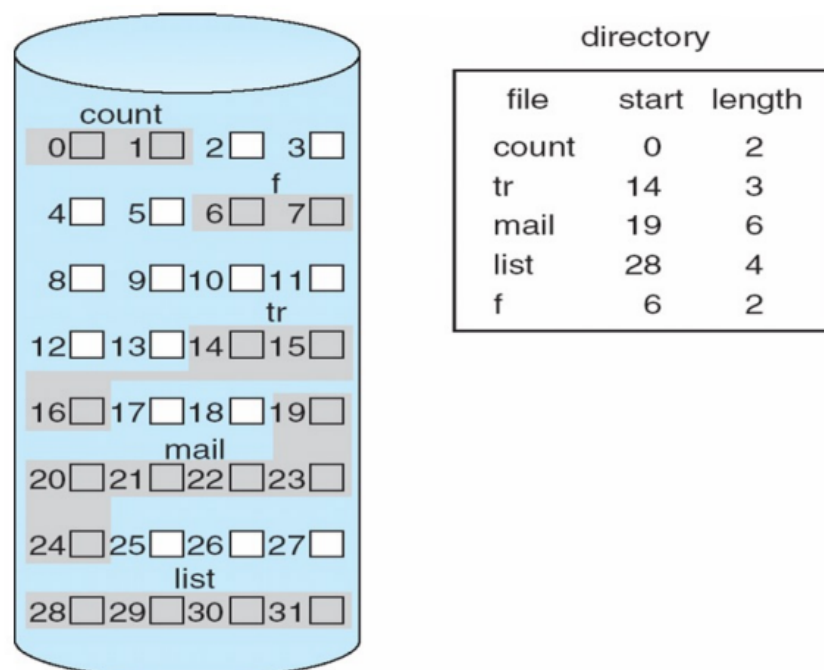
Here is the FCB structure I used. Note that this does not create extra memory space, we just turn the specific portion of volume, originally `uchar*` to `FCB*` for better information storage and retrieval. I have tested that the `sizeof(FCB)` is 32, which is exactly the size of desired FCB. The attributes are self-explanatory.

```
// **32 bytes** File control block. We will turn the FCB part in volume into (FCB*)
// sizeof(FCB) is 32
struct FCB {
    char filename[20]; // maximum size of filename is 20 bytes
    u32 size; // the size of the file **in bytes**
    u16 modified_time; // the last modified time
    u16 creation_time;
    u16 start_block_idx; // the index of the first of its contiguous blocks
    bool is_on;
};
```

Allocation strategy


The maximum file size allowed is the total content size of files, 1024 KB. I use dynamic, contiguous allocation together with compaction algorithm to maintain the FS. The dynamic scheme is to allow storing this maximum size of file.

The contiguous allocation is one such that adjacent blocks store the file content sequentially, as illustrated in the figure:



The compaction algorithm is utilized when there is fragmentation and a newly written file cannot file enough space. In my compaction algorithm implementation, I maintain a pointer to the first unused block and the first used block moving forward together. They

are constantly swapped, which in effect “compacts” all the used blocks to the front. In the mean time, we update the FCB’s content block attribute.



Compaction algorithm

For the superblock, I use the bitmap, which uses one bit for one content block to indicate on or off.

Designing the APIs

All the required APIs `fs_open`, `fs_read`, `fs_write`, `fs_gsys` (including `LS_S`, `LS_D`, `rm` operations) are implemented. The `fs_open` returns an `fp`, which is the index of the FCB in the FCB array. Another interesting point is for the `LS_D` and `LS_S` operations. I did not use external storage for these two sorting operations. Instead, my implementation is simple, which is in each time, traverse all files and find the largest element that is not printed. This does not need to re-place the blocks. The following figure shows my implementation of `LS_D`.

```

printf("===sort by modified time===\n");

int last_item_time = (1<<15); // trace the time of last printed file
// print the most recent modified file before the last item
for (int i = 0; i < file_count; i++)
{
    int latest_modified_time = 0;
    FCB latest_fcb;
    for (int j = 0; j < fs->FCB_ENTRIES; j++)
    {
        FCB fcb = START_OF_FCB[j];
        if (fcb.is_on && (fcb.modified_time > latest_modified_time) && (fcb.modified_time < last_item_time))
        {
            latest_fcb = fcb;
            latest_modified_time = fcb.modified_time;
        }
    }
    last_item_time = latest_fcb.modified_time;
    printf("%s\n", latest_fcb.filename);
    // printf("%s  time%d\n", latest_fcb.filename, last_item_time);
}

```

The LS_S is more tricky but still uses the above idea. We do three traverses in total. First we traverse each item to get the largest unprinted size of files. Then we traverse to get the count of the largest size files. Then find the file with the file size of largest_file_size and the **earliest created time** among all unprinted items.

Bonus task design

The bonus task is based upon the basic task with modification to add files for directories.

Firstly, we need to add a new attribute to trace the current working directory. I add to the `fs` struct.


```

struct FileSystem {
    uchar *volume;
    int SUPERBLOCK_SIZE;
    int FCB_SIZE;
    int FCB_ENTRIES;
    int STORAGE_SIZE;
    int STORAGE_BLOCK_SIZE;
    int MAX_FILENAME_SIZE;
    int MAX_FILE_NUM;
    int MAX_FILE_SIZE;
    int FILE_BASE_ADDRESS;

    int cwd;    // current working directory's fcb index, **not** block index
};

```

Then, the FCB should be different to record each file's directory index.

We squeeze the `is_on` and `is_dir` 1-bit attribute to the first two bits of `u32 size`. So this is 32 bytes again.

```

// File control block
// this is 32 bytes. We turn the FCB portion in the volume into (FCB*). So no extra space
struct FCB {
    char filename[20]; // maximum size of filename is 20 bytes
    u32 size; // the size of the file **in bytes**
    u16 modified_time; // the last modified time
    u16 creation_time;
    u16 start_block_idx; // the index of the first of its contiguous blocks

    // if the FCB is a directory, it stores its parent dir index.
    // if the FCB is root directory, its parent dir index will be -1
    int16_t dir_idx; // the index of the directory that the file is in; for a dir, it is its own idx
};

```

My implementation does not use extra global memory. Other implementations are similar to basic task. Because we record the file contents or subdirectory names in the content of directory, we need to traverse and check the filename match and the `dir_idx` match the `cwd`.

All required operations are supported, including the extra command `MKDIR`, `PWD`, `CD`, `RM_RF` and `CD_P`, in addition to the ones in basic task.

My implementation also supports absolute addressing to increase robustness.

Project reflection and conclusion

Several problems I met in this assignment gave me valuable experience in solving them. The first is

about data structure used to implement FCB in bonus. At first I thought 32 bytes is not enough and want to implement a doubly linked list in the bonus. However, later I found that I could squeeze the FCB on/off bits, and can use a filename match traversing strategy instead of linked list. This allows more efficient storage utilization.

I think this project is a valuable experience for learning the FS, including dynamic allocation, contiguous allocation, compaction. I also learn the technique of writing CUDA programs, which are somewhat like C/C++ but have restricted access to some standard library routines, like `memcpy`.